

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

Searching. Sorting.

Lect. PhD. Arthur Molnar

Babes-Bolyai University

Overview

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

1 Searching

- The searching problem
- Searching with unordered keys
- Searching with ordered keys

2 Sorting

- The sorting problem
- (Binary) Insertion sort
- Quick Sort
- Merge Sort
- TimSort

Let's define the search problem

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching problem

Searching with
unordered keys

Searching with
ordered keys

Sorting

The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort

- Data are available in the internal memory, as a sequence of records (k_1, k_2, \dots, k_n)
- We search for the record that has a certain value for one of its fields, called the **search key**.
- We return the position of the record in the given sequence, or -1 if the **search key** was not found
- Depending on the structure of the sequence of records, we distinguish two possibilities:
 - Searching with unordered keys
 - Searching with ordered keys

Searching with unordered keys

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Searching with
ordered keys

Sorting

The sorting
problem

(Binary)

Insertion sort

Quick Sort

Merge Sort

TimSort

Problem specification

- **Data:** A is a sequence of records with n elements, where $n \in \mathbb{N}$.
- **Result:** $index$, where $(0 \leq index \leq n - 1)$, the search key's position in sequence A , or $index = -1$, if the key was not found.

Searching with unordered keys

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Searching with
ordered keys

Sorting

The sorting
problem

(Binary)

Insertion sort

Quick Sort

Merge Sort

TimSort

Sequential search (iterative and recursive)

Examine the source code in **ex08_sequential_search.py**

What's the computational complexity:

- Best case: search key is the first element checked, so $T(n) = 1 \in \Theta(1)$
- Worst case: search key is not found, so $T(n) = n \in \Theta(n)$

Searching with unordered keys

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

What's the computational complexity:

- Average case: let's assume¹ 50% of searches end with the element not found, and when the element is found, it has the same probability of being on any position in the list. This time, we simulate $2 * n$ runs; during n runs, the element is not found, and during the remaining runs, it is found on each of the list's n positions;

$$T(n) = \overbrace{\frac{1 + 2 + \dots + n}{2 * n}}^{\text{key found, } n \text{ searches}} + \overbrace{\frac{n + n + \dots + n}{2 * n}}^{\text{key not found, } n \text{ searches}} \in \Theta(n)$$

¹an assumption so that we may provide probability $P(I)$

Searching with unordered keys

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

Sequential search (iterative and recursive)

Examine the source code in `ex08_sequential_search.py`

What about a recursive implementation²?

$$\begin{cases} T(n) = 1, n = 1 \\ T(n) = T(n-1) + 1, n \geq 1 \end{cases}$$

Leads to a telescopic sum that resolves to $T(n) \in \Theta(n)$

²Only for fun ☺, as the stack usually frowns upon this implementation 🔍🔍🔍

Searching with ordered keys

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

Problem specification

- **Data:** A is a sequence of records with n elements, where $n \in \mathbb{N}$ so that $\forall 0 \leq i < j \leq n - 1, A[i] \leq A[j]$
- **Result:** $index$, where $(0 \leq index \leq n - 1)$, the search key's position in sequence A , or $index = -1$, if the key was not found.

Searching with ordered keys

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

Binary search (iterative and recursive)

Examine the source code in **ex09_binary_search.py**. Also pay attention to the *test_binary_search* method.

What's the computational complexity:

- Best case: search key is the first element checked, so $T(n) = 1 \in \Theta(1)$
- Worst case: search key is not found.
$$\begin{cases} T(n) = 1, n = 1 \\ T(n) = T(n/2) + 1, n \geq 1 \end{cases}, \text{ each time halving the}$$

array in $O(1)$ time. Since we can do this $\log_2 n$ times, the complexity in the average and worst cases is $\Theta(\log_2 n)$

The sorting problem

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

Sorting

Rearrange a data collection in such a way that the elements of the collection verify a given order.

Problem specification

- **Data:** A is a sequence of records with n elements, where $A = (a_1, a_2, \dots, a_n)$, $a_i \in \mathbb{R}$, $i = 1, n$
- **Results:** A' , where A' is a permutation of A , having sorted elements: $a'_1 \leq a'_2 \leq \dots \leq a'_n$.

The sorting problem

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

Terminology:

- **Internal sort** - data to be sorted are available in the internal memory
- **Streaming sort** - data becomes available sequentially (being read from the network, a database etc.)
- **In-place sort** - transforms the input data into the output, only using a small additional space. Its opposite is called out-of-place.
- **Sorting stability** - the original order of multiple records having the same key is preserved
- **Adaptive sort** - an algorithm that takes data structure into account while sorting
- **Hybrid sort** - an algorithm that uses two or more sorting algorithms to speed up sorting

Demo

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

Stable sort example

Examine the source code in **ex10_stable_sort.py**

Sorting algorithms

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

A few algorithms that we will study:

- (Binary) Insertion sort
- Quick Sort
- Merge sort
- TimSort

Insertion sort

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

- Traverse the list once, examining each element a single time
- Insert the examined element at the right position in the subsequence of already sorted elements.
- The sub-sequence containing the already processed elements is kept sorted, so that, at the end of the traversal, the entire sequence is sorted.
- Very good for short lists, or lists that are nearly sorted (with only a few elements out of place)

Insertion sort – fact sheet

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

In-place	✓	Stable	✓
Streaming	✓	Space	$O(1)$
Time	Best $O(n)$	Average $O(n^2)$	Worst $O(n^2)$

Binary Insertion sort

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

- An optimization of insertion sort
- Since the sequence of already processed elements is sorted, we use binary search to find the place of each new element

Binary Insertion sort

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

Question

What happens if use (binary) insertion sort to sort **linked lists** instead of **arrays**? What are the differences?

Binary) Insertion sort

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

Example

Examine the source code in **ex11_insertion_sort.py**

Binary Insertion sort – fact sheet

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

In-place	✓	Stable	✓
Streaming	✓	Space	$O(1)$
	Best	Average	Worst
Time	$O(n * \log_2 n)$	$O(n^2)$	$O(n^2)$

Quick Sort

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

Based on the *divide and conquer* technique

- 1 Divide:** partition array into 2 sub-arrays such that elements in the lower part \leq elements in the higher part.

Partitioning

Re-arrange the elements so that the element called pivot occupies the final position in the sub-sequence. If i is that position: $k_j \leq k_i \leq k_l$, for $Left \leq j < i < l \leq Right$

- 2 Conquer:** recursively sort the 2 sub-arrays.
- 3 Combine:** trivial since sorting is done in place.

Quick Sort - time complexity

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

- An example of in-place, divide & conquer, non-stable sort 😊
- Its time complexity depends on the distribution of splits, namely pivot selection. Some strategies:
 - Choose the first or last element
 - Choose a random element, or the one in the middle
 - Median of three elements in the partition (e.g., first, midpoint, last)

Quick Sort - best partitioning

Lecture 03

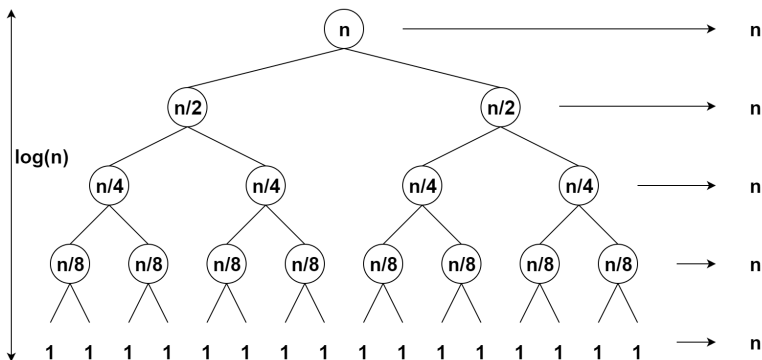
Lect. PhD.
Arthur Molnar

Searching

The searching problem
Searching with unordered keys
Searching with ordered keys

Sorting

The sorting problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort



- We partition n elements $\log_2 n$ times, so
 $T(n) \in \Theta(n \log_2 n)$

Quick Sort - worst partitioning

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

- In the worst case, function Partition splits the array such that one side of the partition has only one element:

$$T(n) = T(1) + T(n-1) + \Theta(n) = T(n-1) + \Theta(n) =$$
$$\sum_{k=1}^n \Theta(k) \in \Theta(n^2)$$

Quick Sort - Worst case

Lecture 03


Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort



- Worst case appears when the input array is sorted or reverse sorted, and we select the first or last element as the pivot
- The n elements are partitioned n times, $T(n) \in \Theta(n^2)$

Quick Sort – fact sheet

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

In-place	✓	Stable	X
Streaming	X	Space	$O(\log_2 n)$
	Best	Average	Worst
Time	$O(n * \log_2 n)$	$O(n * \log_2 n)$	$O(n^2)$

Merge Sort

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

- Based on the divide & conquer approach.
- The list to be sorted is divided in two sub-lists that are sorted separately. The sorted sub-lists are then merged.
- Each sub-list is sorted using the same approach until we get to sub-lists of length 1, which we know are sorted.

Merge Sort

Lecture 03

Lect. PhD.
Arthur Molnar


Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

- + Merge sort has good time complexity and it's easy to write an implementation that works well...
- It has non-constant extra-space complexity
- Time is wasted when merging very short lists (length 1, 2, 4, ...)



Merge Sort – fact sheet

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

In-place	X	Stable	X
Streaming	X	Space	$O(n)$
	Best	Average	Worst
Time	$O(n * \log_2 n)$	$O(n * \log_2 n)$	$O(n * \log_2 n)$

Sorting Algorithm Optimization

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

- There's no really "best" sorting algorithms, as data size, structure, cost of comparisons, CPU cache sizes and other things influence actual run times
- e.g., merge-insertion sort (a.k.a Ford-Johnson algorithm) was known for a long time as the sorting algorithm with the smallest number of element comparisons
- Combine the strengths of several algorithms into **hybrid algorithms** – merge or quick sort that default to insertion sort for short lists

Sorting Algorithm Optimization

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

Sorting algorithms visualized

<https://www.toptal.com/developers/sorting-algorithms>

TimSort

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

- Developed for Python in 2002 by Tim Peters
- Adapted in Java 7, Android, Swift, V8 and Rust
- A hybrid algorithm built on merge sort that employs binary insertion sort for small sections of the list, leverages natural element ordering and in-cache data 😊

A Word of Warning

You'll find **many** resources on the interwebs showing "TimSort", but most of them are either incomplete, or straightforward combinations of merge and insert sort. Some of them are even buggy 😞

The Real TimSort

- The definitive description of the algorithm by its author (<https://svn.python.org/projects/python/trunk/Objects/listsort.txt>) together with its source code (<https://github.com/python/cpython/blob/main/Objects/listobject.c>)
- Some more details regarding it (<https://mail.python.org/pipermail/python-dev/2002-July/026837.html>)

A few other resources for those interested

- Good (video) descriptions (<https://www.awesomealgorithms.com/home/tim-sort>) and (<https://ericmervin.medium.com/what-is-timsort-76173b49bd16>)
- Detailed, step by step description **with Python source code** (<https://vladris.com/blog/2021/12/30/timsort.html>)
- Source code in Python from the RPython project (<https://github.com/reingart/pypy/blob/master/rpython/rlib/listsort.py>)

TimSort

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

In the author's own words...

This describes an adaptive, stable, natural mergesort, modestly called timsort (hey, I earned it ☺). It has supernatural performance on many kinds of partially ordered arrays (less than $\lg(N!)$ comparisons needed, and as few as $N-1$), yet as fast as Python's previous highly tuned samplesort hybrid on random arrays.

In a nutshell, the main routine marches over the array once, left to right, alternately identifying the next run, then merging it into the previous runs "intelligently". Everything else is complication for speed, and some hard-won measure of memory efficiency.

<https://svn.python.org/projects/python/trunk/Objects/listsort.txt>

TimSort

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

Most existing lists include sublists that are already sorted (e.g., [1, 4, 5, 2, 99, 6, 9, 11, 10, 90] has [1, 4, 5] and [6, 9, 11] length-3 sublists)

- 1 Determine the value for *minrun*, a parameter that balances good insert sort performance ($32 < \textit{minrun} < 64$) so that the list is divided into a power of 2 runs
- 2 Sort consecutive runs of length *minrun* using binary insertion sort
- 3 Place runs on a stack ($O(n)$ space complexity)
- 4 Merge consecutive runs from the stack using merge sort (this leads to a stable sort)

TimSort - optimizations

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

- While elements are sorted, keep going even after *minrun*
- A run sorted in descending order can be reversed to ascending in $O(n)$ (check order of first two elements)
- Only call insertion sort for those parts of the run that are not already sorted
- Merging often (and soon after runs are completed) raises the chance that data is (still) in the CPU cache
- Introduce a stack invariant on the length of the merges to ensure balanced merges
- Optimize merging by finding which elements remain in the existing order (e.g., best case when the first element of run A is larger than the last of run B)
- When merging uses data from the same run repeatedly, use an exponential search (galloping) to identify how many more elements can be used

TimSort – fact sheet

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem
Searching with
unordered keys
Searching with
ordered keys

Sorting

The sorting
problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

In-place	✓	Stable	✓
Streaming	X	Space	$O(n)$
	Best	Average	Worst
Time	$O(n)$	$O(n * \log_2 n)$	$O(n * \log_2 n)$